HARIRAMAKRRISHNAN RAMACHANDRAN

+353 89 970 6156 • hariramakrrish@gmail.com • Dublin, Ireland • **Stamp 1G — Full-Time Work Eligible**

PROFESSIONAL EXPERIENCE — 3.5 YEARS

Software Engineer — HCL Technologies, Chennai, India Sep 2021 – Jan 2025

- Architected and deployed **Python**-based feature pipelines on **AWS Lambda** and **S3**, processing approximately **3,500 events/sec** during peak hours to replace a legacy batch job.
- Developed and evaluated **predictive models** using **scikit-learn** on a **~280k-row** dataset for an internal anti-abuse workflow, achieving an **AUC of ~0.86**.
- Designed and implemented **REST APIs** in **Java Spring Boot** for inter-service data exchange, incorporating request validation and **circuit-breaker patterns**.
- Integrated foundational **LLMs** and **embeddings** into agentic AI solutions, facilitating data augmentation, summarization, and contextual reasoning for enterprise client use cases.
- Communicated complex analytical results and model performance to business stakeholders through interactive **Power BI** visualizations and formal presentations.
- Collaborated with engineering teams to develop and refine **data pipelines** using **Spark**, ensuring efficient data wrangling and preparation for model training on cloud platforms.
- Containerized **feature pipeline services** with **Docker** and deployed to an **Kubernetes** cluster (EKS), configuring Helm charts and resource limits for predictable rollouts.

SKILLS

AI/ML – Generative AI, Predictive Analytics, LLMs, Embeddings, Agentic AI, Model Development, Model Evaluation, Feature Engineering, Python, Scikit-learn, PyTorch, TensorFlow, Azure Databricks

Data Engineering – Data Pipelines, Spark, Azure Data Factory, SQL, Data Wrangling, ETL

Programming Languages – Python, SQL, Java, R

Cloud Platforms – AWS (EC2, S3, Lambda, CloudWatch), Azure

DevOps & Orchestration – Docker, Kubernetes, Jenkins, Git

BI & Visualization – Power BI, Tableau, DAX, Power Query

Frameworks & Libraries – Spring Boot, pandas, NumPy, LangChain, LangGraph

PROJECTS

Customer Onboarding Playbook for ML Feature Store

- Developed a guided onboarding flow for two pilot enterprise customers adopting an internal ML feature store, including SDK examples and debug tooling.
- Built a **Python CLI** to validate customer feature definitions against the platform schema, catching approximately **70%** of schema errors at write time.
- Ran technical workshops with customer data teams to upskill their engineers on integrating with the ML platform and feature store.

Fraud Detection Model on Transaction Data

- Engineered a **feature pipeline** in **Python** and **SQL** to extract behavioural signals from a transaction log of approximately **1.1M records**.
- Trained a **LightGBM classifier** on extracted features, tuning class weights to address a **1:140 fraud imbalance** and achieving approximately **0.91 AUC**.
- Packaged the trained model behind a **FastAPI** inference endpoint, implementing input validation and logging for feature drift, achieving a p95 latency of **180ms**.

EDUCATION

M.Sc. Data Analytics — National College of Ireland, Dublin Feb 2026

B.E. Computer Science — SNS College of Technology, Coimbatore, India 2021

CERTIFICATIONS

- IELTS Academic: Overall Band Score **7 out of 9**, issued by IELTS Official, Test Date: 28/AUG/2023
- **Full Stack: Java Spring Boot** and Angular (Great Learning)
- Certificate of Completion in Web Development from **University of California**, Davis (Coursera)
- Certificate of Completion in Python programming from **University of Michigan** (Coursera)