

HARIRAMAKRRISHNAN RAMACHANDRAN

+353 89 970 6156 • hariramakrrish@gmail.com • Dublin, Ireland • Stamp 1G — Full-Time Work Eligible

PROFESSIONAL EXPERIENCE

Software Engineer — HCL Technologies, Chennai, India

Sep 2021 – Jan 2025

- Designed and implemented **Python**-based data pipelines on **AWS Lambda** and **S3** to process approximately **3,500 events/sec** during peak hours, replacing a less efficient batch job.
- Developed and maintained **machine learning models** using **scikit-learn** on a **~280k-row** dataset for internal anomaly detection workflows, achieving an **AUC of ~0.86** on held-out data.
- Collaborated with enterprise customers as the primary technical contact for onboarding, gathering **feature pipeline** requirements directly with their data science teams and translating them into **Python** implementations.
- Built and validated **dashboards and reports** using **Power BI** to track key business trends and performance indicators, presenting findings to senior stakeholders.
- Applied advanced statistical techniques such as **OLS** and **clustering** to conduct exploratory data analysis on large, complex datasets, translating analytical findings into actionable insights.
- Designed and shipped **REST APIs** in **Java Spring Boot** for cross-service data exchange, incorporating request validation and circuit-breaker patterns for enhanced reliability.
- Containerized data processing services using **Docker** and deployed them to an **AWS EKS** cluster, configuring basic resource limits for predictable performance.

SKILLS

Programming Languages – Python, R, SQL, Java

Data Science & ML – Machine Learning, Deep Learning, Scikit-learn, TensorFlow, PyTorch, OLS, Multinomial Logistic Regression, LDA, Clustering, Segmentation, AI, Generative AI, Agentic AI

Big Data & Cloud – Databricks, Spark, Hadoop, AWS (EC2, S3, Lambda, CloudWatch), Docker, Kubernetes, EKS

BI & Visualization – Power BI, Tableau, DAX, Power Query

Backend Development – Spring Boot, Java, REST APIs, Hibernate, JPA

Databases – PostgreSQL, MySQL, Oracle

PROJECTS

Customer Onboarding Playbook for an ML Feature Store

- Worked directly with two pilot customers to design and ship a guided onboarding flow for an internal ML feature store, including SDK examples and debug tooling that compressed time-to-first-feature from ~6 weeks to 9 days.
- Built a **Python CLI** to validate customer feature definitions against the platform schema before submission, catching ~70% of schema errors at write time instead of at training time.
- Ran two follow-up workshops with customer data teams to upskill their internal engineers on feature pipeline development and platform best practices.

Transaction Analysis and Fraud Detection Model

- Engineered a feature pipeline in **Python** and **SQL** to extract behavioural signals from a transaction log of roughly 1.1M records.
- Trained a **LightGBM classifier** on the resulting features, tuning class weights to handle a 1:140 fraud imbalance and reaching **~0.91 AUC / 0.46 F1** on a held-out month.
- Packaged the model behind a **FastAPI** inference endpoint with input validation, achieving p95 latency around **180ms**.

EDUCATION

M.Sc. Data Analytics — National College of Ireland, Dublin

Feb 2026

CERTIFICATIONS

- IELTS Academic: Overall Band Score **7 out of 9**, issued by IELTS Official, Test Date: 28/AUG/2023
- **Full Stack: Java Spring Boot** and Angular (Great Learning)
- Certificate of Completion in Web Development from **University of California**, Davis (Coursera)
- Certificate of Completion in Python programming from **University of Michigan** (Coursera)